

Introduction to database systems

2021/2022

Homework 2

SÜLEYMAN GÖLBOL [Erasmus Student]

Task 1 (50 %): Linear Hash Index

Given is the following linear hash index. Bucket is split every time when a new overflow page is added.

Show the state of the index after adding each key, 18, 20 and 27.

h_1	h_0	Level = 1, N = 4			
000	00	8*	16*		
001	01	5*	13*	21*	
010	10	2*	10*	34*	70*
011	11	7*	15*	59*	91*
100	00	12*	20*	28*	

Next = 1

Adding 18 (next=1)

	h_1	h_0	Level = 1, N = 4			
0	000	00	8*	16*		
1	001	01	5*	13*	21*	
2	010	10	2*	10*	34*	70*
3	011	11	7*	15*	59*	91*
4	100	00	12*	20*	28*	

Overflow Pages

RULES:

Split condition is when a new overflow page.
Before moving to other level, all 4 pages must be split.

formula for insertion: $\text{key} \bmod (2^{\text{level}} \cdot N)$
formula for redistribution: $\text{key} \bmod (2^{\text{level}+1} \cdot N)$

(level=1) so $18 \bmod (2^1 \cdot 4) = 2$

18* will be added to row number 2

split condition is when a new overflow page

redistribute row=next which is 1

So redistribute;

$5 \bmod (2^1 \cdot 4) = 5$

$13 \bmod (2^1 \cdot 4) = 5$

$21 \bmod (2^1 \cdot 4) = 5$

(18 was added so next++ => next = 2)

The new table

DECIMAL	h_1	h_0	Level = 1, N = 4, Next = 2				Overflow Pages			
0	000	00	8*	16*						
1	001	01								
2	010	10	2*	10*	34*	70*	18*			
3	011	11	7*	15*	59*	91*				
4	100	00	12*	20*	28*					
5	101	01	5*	13*	21*					

Adding 20

$20 \bmod (2^1 \cdot 4) = 4$ but 20 is already added so we don't have to do anything.

Adding 27

$27 \bmod (2^1 \cdot 4) = 3$

Redistribute row=next which is 2

$2 \bmod (2^1 \cdot 4) = 2$

$10 \bmod (2^1 \cdot 4) = 2$

$34 \bmod (2^1 \cdot 4) = 2$

$70 \bmod (2^1 \cdot 4) = 6$

$18 \bmod (2^1 \cdot 4) = 2$

next++ => so next=3

DECIMAL	h_1	h_0	Level = 1, N = 4, Next = 2				Overflow Pages			
0	000	00	8*	16*						
1	001	01								
2	010	10	2*	10*	34*	18*				
3	011	11	7*	15*	59*	91*	27*			
4	100	00	12*	20*	28*					
5	101	01	5*	13*	21*					
6	110	10	70*							

of index entries per page => should be rounded DOWN
 # of within cotain index level => should be rounded UP
 (15,6 = 16 Pages)

Task 2 (50 %): Query evaluation and optimisation

The most important tables in the library information system are the following.

Books (bid, author, title, publisher, year);
 Members (mid, name, surname, address, telephone);
 Rental (rid, mid, bid, eid, date);
 Employee (eid, name, surname, address, telephone)

The following information is given.

Multiple-byte units			
Decimal		Binary	
Value	Metric	Value	IEC
1000	KB	1024	KB

my assumption for 1 KB is 1024 bytes
 (like prominently used by the Microsoft Windows operating system)

1 page on the disk = 8KB

|Books| = 1.000.000 records, 320 bytes, 25 records/page, 40000 pages

|Members| = 10.000 records, 200 bytes, 40 records/page, 250 pages

|Rental| = 300.000 records, 40 bytes, 200 records/page, 1500 pages

|Employee| = 100 records, 200 bytes, 40 records/page, 3 pages

We can have 10000 pages in the buffer.

The following indexes are set and available on the database:

- unclustered B+ index on Rental.date attribute
- clustered B+ index on Books.year attribute and
- hash indexes on all relation keys.

Assume that all tables are ordered by the values of their keys.

Our SUPB has the following join algorithms available:

- Sort-merge Join
- Hash Join

PART 1 (20%). What is the size of B+ index on attribute

Books.year? DO NOT use estimations. Write down all assumptions.

clustered B+ index is available.

B+ tree

Except root, minimum 50 occupancy
 $d \leq m \leq 2d$
 Order of tree

1 page = 8.1024 bytes
 on disk

Books.year = 12 bytes (my assumption for Books.year)
 page index identifier size = 4 bytes (my assumption for page identifier)
 +
 Size of 1 page index entry
 $12 + 4 = 16$ bytes

Books.year = 12 bytes (my assumption for Books.year)
 data record identifier size = 8 bytes (my assumption for record identifier)
 +
 Size of 1 data entry
 $12 + 8 = 20$ bytes

Number of index entries/page = $\frac{8.1024}{16} = 512$

Number of data record entries/page = $\frac{8.1024}{20} = 410$

Leaf (Level 2) $\frac{1000000}{410} \approx 2440$ pages (block)

Level 1 $\frac{2440}{512} = 4,76 \approx 5$ pages

Root (Level 0) $\frac{5}{512} = 0,0097 \approx 1$ page

Total 2446 pages

$2446 \times 8 \text{ KB} = \frac{19568 \text{ KB}}{1024} = 19,10 \text{ MB}$


PART 2 (30%). Translate the following SQL statement into a relational algebra tree (query execution plan) and find a plan that reads the least pages:

SELECT B.author, B.title
FROM Rental R, Books
B
WHERE R.bid=B.bid AND B.year='2000' AND R.date > '1/1/2019'

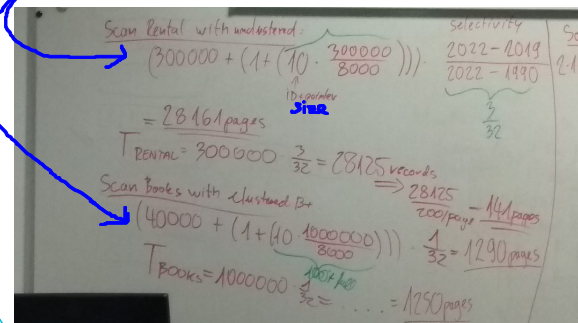
Rental.date ranges from 1.1.1990 to 31.12.2021. Books.year ranges from 1990 to 2021.
For the sake of simplicity, assume that we are already in year 2022.

Costs of access path should be calculated with known formula.

Write down all assumptions.



cost of scan



Scanning Rental

1/1/2019 to 31/12/2021 is 3 years.
Totally there are 32 years.

$$300000 \cdot \frac{3}{32} = 28125 \text{ records}$$

$$200 \frac{\text{records}}{\text{page}}$$

$$= 140,6 \sim 141 \text{ page}$$

scan with unclustered
b+ index and store to
Trental

Scanning Books

From 1990 to 2021
there are 32 years.

$$\frac{1 \text{ million records}}{32} = 31250 \text{ records}$$

$$25 \frac{\text{records}}{\text{page}}$$

$$= 1250 \text{ pages}$$

scan with clustered b+
index and store to
Tbooks

Sort Merge Join:

$$\text{Formula: } 2M(1 + \log_{B-1}(\frac{M}{B})) + 2N(1 + \log_{B-1}(\frac{N}{B})) + M + N$$

$$B=10K$$

$$B-1=999$$

For T_1 Sorting \Rightarrow

$$2 \cdot 141 \cdot (1 + \frac{\log_{999} 141}{\log_{999} 10000}) + 2 \cdot 1250 \cdot (1 + \frac{\log_{999} 1250}{\log_{999} 10000}) + 141 + 1250$$

$$110 \quad 1750 \quad 1391$$

$$= 3232 \text{ pages}$$

Better
option

Hash Join

$$3(M+N)$$

Table1 Table2

$$3(141+1250) = 3.1391$$

$$= 4173 \text{ pages}$$